

Hammad Ali Tahir

[Portfolio](#) | [Github](#) | +923334247884 | hammadalitahir8@gmail.com | [Linkedin](#)

SUMMARY

BS Computer Science fresh graduate specializing in ML, NLP, and LLM fine-tuning, with hands-on experience building RAG pipelines, deploying full-stack AI applications, and achieving production-grade model accuracy. Seeking an ML Engineer, NLP Engineer, or AI Engineer role to build and deploy production-ready AI systems at scale.

SKILLS

Languages: Python, SQL

ML/AI Frameworks: TensorFlow, Keras, PyTorch, Scikit-learn, XGBoost

LLM/NLP: Hugging Face, LangChain, LangGraph, OpenAI API, RAG, Fine-tuning, LoRA, QLoRA, Prompt Engineering, Quantization, NLTK

Vector Databases: Pinecone, Qdrant, ChromaDB

Backend: FastAPI, Flask, Streamlit, RESTful APIs, Redis

DevOps/Cloud: Docker, Git, Render, Vercel, CI/CD

Data: Pandas, NumPy, EDA, Feature Engineering

Concepts: CNNs, NLP, Supervised/Unsupervised Learning, Ensemble Methods, Vector Similarity Search

Soft Skills: Adaptable, fast learner, strong communicator, collaborative team player

EXPERIENCE

[AI Engineer | AI-Based Legal Case Management System](#)

Uraan AI Techathon 1.0

Sep 2025 - Oct 2025

- Built a production-grade AI LegalTech SaaS platform with intelligent case classification, automated triage, and semantic precedent retrieval using ML, NLP, and RAG pipelines.
- Achieved **91.5%** classification accuracy and **95.2%** prioritization precision using ensemble models (SVM, Naïve Bayes, Random Forest), improving legal data analysis efficiency by **40%+**.
- Developed full-stack AI application (React.js, Python/Streamlit) integrating Hugging Face transformer embeddings, Qdrant vector database, and Llama-3.1-8B LLM for context-aware legal research and generative AI insights.
- Deployed cloud-native solution on Render and Vercel with CI/CD pipelines, RESTful APIs, and vector similarity search for scalable Legal AI delivery.

[AI Developer | Academic Chatbot System](#)

University of Education Lahore

Jan 2026 - Apr 2026

- Designed and deployed a full-stack AI academic assistant for University of Education Lahore using FastAPI, React/Vite, OpenAI GPT-4o-mini, Redis memory, and Pinecone-based RAG across 4 university knowledge namespaces, achieving **92% answer relevancy** on evaluated academic queries.
- Developed an advanced Smart RAG pipeline incorporating query enhancement, namespace routing, **metadata filtering**, cross-encoder reranking, chunk relevance grading, query rewriting, query decomposition, and hallucination detection to generate grounded responses with source attribution and relevance scoring.
- Built a production-ready ChatGPT-style Progressive Web Application (PWA) featuring Server-Sent Events (SSE) streaming, Zustand state management, a mobile-first interface, configurable RAG controls, feedback collection, and LangSmith observability, with **100+ production traces logged** for retrieval and debugging evaluation.

PROJECTS

[Loan Approval Prediction System](#)

Python | Flask | Scikit-learn | Docker | Render

Jun 2025 - Jul 2025

- Built an end-to-end Loan Approval Prediction System using Python, Flask, scikit-learn, and Gunicorn, implementing a complete ML pipeline, data ingestion, feature engineering (StandardScaler, OneHotEncoder, ColumnTransformer), and automated model selection across 6 classifiers (Gradient Boosting, Random Forest, AdaBoost, SVC, Logistic Regression, Decision Tree) achieving 97.6% accuracy with **0.99 ROC-AUC**.
- Deployed a production-ready ML web application on Render with Docker support, featuring RESTful prediction APIs, modular training/prediction pipelines, health monitoring, structured logging, and custom exception handling, following industry-standard CI/CD and scalable software engineering practices.

[GitHub Research Agent](#)

Python | FastAPI | LangGraph | Render | Vercel

Oct 2025 - Nov 2025

- Built an AI-powered GitHub research agent using LangGraph and FastAPI, implementing an automated multi-step pipeline that retrieves repository metadata (stars, forks, languages, activity), analyzes architecture, inspects issues, and synthesizes caching-enabled research reports via a **ReAct agent** equipped with live GitHub tools.
- Developed a full-stack application with a React + Vite frontend deployed on Vercel and a Python/FastAPI backend deployed on Render, featuring a responsive real-time chat interface for interactive follow-up Q&A powered by a LangGraph ReAct reasoning loop.

EDUCATION

BS Computer Science-University of Education Lahore

- *Relevant Coursework: Linear Algebra, Statistics, Calculus-I, OOP, DSA, Machine Learning, Data Science, Database Systems*

CERTIFICATION

- **Supervised Machine Learning: Regression and Classification** — [DeepLearning.AI](#) & Stanford University | 2025
- **Oracle Cloud Infrastructure 2025 Certified AI Foundations Associate** — Oracle University | Aug 2025